

Tree search:

- A* is optimal if heuristic is admissible
- UCS is a special case ($h = 0$)

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

Graph search:

- A* optimal if heuristic is consistent
- UCS optimal ($h = 0$ is consistent)

- Admissibility: heuristic cost \leq actual cost to goal

$$h(A) \leq \text{actual cost from A to G}$$

- Consistency: heuristic "arc" cost \leq actual cost for each arc

$$h(A) - h(C) \leq \text{cost}(A \text{ to } C)$$

Consistency implies admissibility

- The f value along a path never decreases

$$h(A) \leq \text{cost}(A \text{ to } C) + h(C)$$

Strategies: DFS: stack BFS: queue UCS: cumulative cost

α : MAX's best option on path to root
 β : MIN's best option on path to root

```
def max-value(state,  $\alpha$ ,  $\beta$ ):
    initialize  $v = -\infty$ 
    for each successor of state:
         $v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$ 
    if  $v \geq \beta$  return  $v$ 
     $\alpha = \max(\alpha, v)$ 
    return  $v$ 
```

```
def min-value(state,  $\alpha$ ,  $\beta$ ):
    initialize  $v = +\infty$ 
    for each successor of state:
         $v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$ 
    if  $v \leq \alpha$  return  $v$ 
     $\beta = \min(\beta, v)$ 
    return  $v$ 
```

Variables: WA, NT, Q, NSW, V, SA, T

Domains: $D = \{\text{red, green, blue}\}$

Constraints: adjacent regions must have different colors

Implicit: $WA \neq NT$

Explicit: $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), \dots\}$

Solutions are assignments satisfying all constraints, e.g.:

$\{WA=\text{red}, NT=\text{green}, Q=\text{red}, NSW=\text{green}, V=\text{red}, SA=\text{blue}, T=\text{green}\}$

Discrete Variables

- Finite domains
 - Size d means $O(d^n)$ complete assignments
 - E.g., Boolean CSPs, including Boolean satisfiability (NP-complete)
- Infinite domains (integers, strings, etc.)
 - E.g., job scheduling, variables are start/end times for each job
 - Linear constraints solvable, nonlinear undecidable

Continuous variables

- Linear constraints solvable in polynomial time by Linear Programming methods
- Varieties of Constraints
 - Unary constraints involve a single variable (equivalent to reducing domains), e.g.: $SA \neq \text{green}$
 - Binary constraints involve pairs of variables, e.g.: $SA \neq WA$
 - Higher-order constraints involve 3 or more variables: e.g., cryptarithmic column constraints

Backtracking Search

Filtering: Keep track of domains for unassigned variables and cross off bad options.

Forward checking: Cross off values that violate a constraint when added to the existing assignment.

```

function BACKTRACKING-SEARCH(csp) returns solution/failure
    return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
    if assignment is complete then return assignment
    var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
        if value is consistent with assignment given CONSTRAINTS[csp] then
            add { var = value } to assignment
            result ← RECURSIVE-BACKTRACKING(assignment, csp)
            if result ≠ failure then return result
            remove { var = value } from assignment
    return failure

```

Arc Consistency

An arc $X \rightarrow Y$ is **consistent** iff for every x in the tail there is *some* y in the head which could be assigned without violating a constraint

Remember: Delete from the tail!

- A simple form of propagation makes sure **all** arcs are consistent:
- Important: If X loses a value, neighbors of X need to be rechecked!
- Arc consistency detects failure earlier than forward checking

MRV

- Variable Ordering: Minimum remaining values (MRV):
 - Choose the variable with the fewest legal left values in its domain

Tree-Structured CSPs

- Remove backward: For $i = n : 2$, apply RemoveInconsistent(Parent(X_i), X_i)
- Assign forward: For $i = 1 : n$, assign X_i consistently with Parent(X_i)
- Claim1: After backward pass, all root-to-leaf arcs are consistent
- Claim2: If root-to-leaf arcs are consistent, forward assignment will not backtrack

Linear Classifiers: Perceptrons

- Inputs are **feature values**
- Each feature has a **weight**
- Sum is the **activation**

$$\text{activation}_w(x) = \sum w_i \cdot f_i(x) = w \cdot f(x)$$

- If the activation is:
 - Positive, output +1
 - Negative, output -1

Binary Perceptron

- Start with weights = 0
- For each training instance:
 - Classify with current weights

$$y = \begin{cases} +1 & \text{if } w \cdot f(x) \geq 0 \\ -1 & \text{if } w \cdot f(x) < 0 \end{cases}$$

- If correct (i.e., $y=y^*$), no change!
- If wrong: adjust the weight vector by adding or subtracting the feature vector. Subtract if y^* is -1.

$$w = w + y^* \cdot f$$

Multiclass Perceptron

- Start with weights = 0
- For each training instance:
 - Start with all weights = 0
 - Pick up training examples one by one
 - Predict with current weights

$$y = \arg \max_y w_y \cdot f(x)$$

- If correct, no change!
- If wrong: lower score of wrong answer, raise score of right answer

$$w_y = w_y - f(x)$$

$$w_{y^*} = w_{y^*} + f(x)$$

Markov Decision Processes (MDP)

- An MDP is defined by:
 - A set of states $s \in S$
 - A set of actions $a \in A$
 - A transition function $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - Also called the model or the dynamics
 - A reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
 - A start state
 - Maybe a terminal state
- MDPs are non-deterministic search problems
 - One way to solve them is with expectimax search

The value (utility) of a state s :

$V^*(s)$ = expected utility starting in s and acting optimally

The value (utility) of a q-state (s,a) :

$Q^*(s,a)$ = expected utility starting out having taken action a from state s and (thereafter) acting optimally

The optimal policy:

$\pi^*(s)$ = optimal action from state s

Bellman Equations:

$$V^*(s) = \max_a Q^*(s, a)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

Value Iteration:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

Policy Iteration:

$$V_0^\pi(s) = 0$$

$$V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$$

$$\pi_{i+1}(s) = \arg \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k^{\pi_i}(s')]$$

Sample-Based Policy Evaluation:

$$sample_1 = R(s, \pi(s), s'_1) + \gamma V_k^\pi(s'_1)$$

$$sample_2 = R(s, \pi(s), s'_2) + \gamma V_k^\pi(s'_2)$$

...

$$sample_n = R(s, \pi(s), s'_n) + \gamma V_k^\pi(s'_n)$$

$$V_{k+1}^\pi(s) \leftarrow \frac{1}{n} \sum_i sample_i$$

Known MDP: Offline Solution

Goal	Technique
Compute V^*, Q^*, π^*	Value / policy iteration
Evaluate a fixed policy π	Policy evaluation

Unknown MDP: Model-Based

Goal	Technique
Compute V^*, Q^*, π^*	VI/PI on approx. MDP
Evaluate a fixed policy π	Approx. MDP

Unknown MDP: Model-Free

Goal	Technique
Compute V^*, Q^*, π^*	Q-learning
Evaluate a fixed policy π	Value Learning

Temporal Difference Learning:

Sample of $V(s)$: $sample = R(s, \pi(s), s') + \gamma V^\pi(s')$

Update to $V(s)$: $V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + (\alpha)sample$

Same update: $V^\pi(s) \leftarrow V^\pi(s) + \alpha(sample - V^\pi(s))$

Q-Learning:

$sample = R(s, a, s') + \gamma \max_{a'} Q(s', a')$

$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha)[sample]$

$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$

SARSA:

$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$

Probability Distributions:

$$\forall x \ P(X = x) \geq 0 \quad \text{and} \quad \sum_x P(X = x) = 1$$

Joint Distributions:

▪ Must obey: $P(x_1, x_2, \dots, x_n) \geq 0$

$$\sum_{(x_1, x_2, \dots, x_n)} P(x_1, x_2, \dots, x_n) = 1$$

Normalization Trick:

1. **SELECT** the joint probabilities matching the evidence.
2. **NORMALIZE** the selections (make them sum to 1)

▪ Conditional probability

$$P(x|y) = \frac{P(x, y)}{P(y)}$$

▪ Product rule

$$P(x, y) = P(x|y)P(y)$$

▪ Chain rule


$$\begin{aligned} P(X_1, X_2, \dots, X_n) &= P(X_1)P(X_2|X_1)P(X_3|X_1, X_2) \dots \\ &= \prod_{i=1}^n P(X_i|X_1, \dots, X_{i-1}) \end{aligned}$$

▪ X, Y independent if and only if: $\forall x, y : P(x, y) = P(x)P(y)$

D-Separation:

Active Triples

Inactive Triples



1. Prior Sampling:

全部采样，然后算每个出现的概率

2. Rejection Sampling:

采样时把不满足 evidence 的样本舍弃，然后计算每个出现的概率

3. Likelihood Weighting (从上到下)

▪ if X_i is an evidence variable

- $X_i = \text{observation } x_i \text{ for } X_i$
- Set $w = w * P(x_i | \text{Parents}(X_i))$

▪ else

- Sample x_i from $P(X_i | \text{Parents}(X_i))$

4. Gibbs Sampling (上下都可)

- Choose a non-evidence variable X
- Resample X from $P(X | \text{all other variables})$

$$P(S|+c, +r, -w) = \frac{P(S, +c, +r, -w)}{P(+c, +r, -w)}$$

Markov Chain: Stationary Distributions

$$P_{\infty}(X) = P_{\infty+1}(X) = \sum_x P(X|x)P_{\infty}(x)$$

Page Rank:

$$PR_{t+1}(P_i) = \sum_{P_j} \frac{PR_t(P_j)}{C(P_j)}, \text{ 其中 } PR_t(P_j) \text{ 初始为 } 1/\text{结点数}, C(P_j) \text{ 为结点出度}$$

Transition Matrix: P_{ij} 表示第 i 行第 j 列为第 i 个结点转移到第 j 个元素的概率

Hidden Markov Models: The Forward Algorithm

$$P(x_t|e_{1:t}) \propto P(x_t, e_{1:t})$$

$$= P(e_t | x_t) \sum_{x_{t-1}} P(x_t | x_{t-1}) P(x_{t-1}, e_{1:t-1})$$

$$P(x_t|e_{1:t}) \propto P(x_t, e_{1:t})$$

$$= \sum_{x_{t-1}} P(e_t | x_t, x_{t-1}) P(x_t | x_{t-1}) P(x_{t-1}, e_{1:t-1})$$

$$P(x_t|e_{1:t}) \propto P(x_t, e_{1:t})$$

$$= \sum_{x_{t-1}} P(e_t | x_t, e_{t-1}) P(x_t | x_{t-1}) P(x_{t-1}, e_{1:t-1})$$

$$P(x_t, x_{t+1}|e_{1:t}) \propto P(x_t, x_{t+1}, e_{1:t})$$

$$= \sum_{x_{t-1}} P(e_t | x_t) P(x_{t+1} | x_t, x_{t-1}) P(x_t, x_{t-1}, e_{1:t-1})$$

$$P(X_2)$$

$$P(X_1|e_1)$$

$$\begin{aligned} P(x_2) &= \sum_{x_1} P(x_1, x_2) \\ &= \sum_{x_1} P(x_1) P(x_2|x_1) \end{aligned}$$

$$\begin{aligned} P(x_1|e_1) &= P(x_1, e_1)/P(e_1) \\ &\propto_{X_1} P(x_1, e_1) \\ &= P(x_1) P(e_1|x_1) \end{aligned}$$

Particle Filtering

- Our representation of $P(X)$ is now a list of N particles (samples)
- Each particle is moved by sampling its next position from the transition model

$$x' = \text{sample}(P(X'|x))$$

- Similar to likelihood weighting, downweight samples based on the evidence

$$w(x) = P(e|x)$$

$$B(X) \propto P(e|X)B'(X)$$

- Rather than tracking weighted samples, we resample;
- N times, we choose from our weighted sample distribution (draw with replacement);
- This is equivalent to renormalizing the distribution;
- Now the update is complete for this time step, continue with the next one.

Naïve Bayes

$|Y|$ parameters

$$P(Y, F_1 \dots F_n) = P(Y) \prod_i P(F_i|Y)$$

$|Y| \times |F|^n$ values

$n \times |F| \times |Y|$
parameters

$$P(Y, f_1 \dots f_n) = \begin{bmatrix} P(y_1, f_1 \dots f_n) \\ P(y_2, f_1 \dots f_n) \\ \vdots \\ P(y_k, f_1 \dots f_n) \end{bmatrix} \Rightarrow \frac{\begin{bmatrix} P(y_1) \prod_i P(f_i|y_1) \\ P(y_2) \prod_i P(f_i|y_2) \\ \vdots \\ P(y_k) \prod_i P(f_i|y_k) \end{bmatrix}}{P(f_1 \dots f_n)} \Rightarrow P(Y|f_1 \dots f_n)$$

$$\text{prediction}(f_1, \dots, f_n) = \underset{y}{\operatorname{argmax}} P(Y = y) \prod_{i=1}^n P(F_i = f_i | Y = y)$$

Laplace Smoothing

$$P_{LAP,k}(x) = \frac{c(x) + k}{N + k|X|}$$

$$P_{LAP,k}(x|y) = \frac{c(x, y) + k}{c(y) + k|X|}$$

